

```
import pandas as pd
import numpy as np # Import the numpy library
url = "https://raw.githubusercontent.com/plotly/datasets/master/diabetes.csv"
df = pd.read_csv(url)
```

```
df.info()
df.describe()
df.isnull().sum()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Pregnancies                          768 non-null    int64
 1   Glucose                              768 non-null    int64
 2   BloodPressure                        768 non-null    int64
 3   SkinThickness                       768 non-null    int64
 4   Insulin                             768 non-null    int64
 5   BMI                                  768 non-null    float64
 6   DiabetesPedigreeFunction             768 non-null    float64
 7   Age                                  768 non-null    int64
 8   Outcome                              768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
0
Pregnancies    0
Glucose         0
BloodPressure   0
SkinThickness   0
Insulin         0
BMI             0
DiabetesPedigreeFunction  0
Age             0
Outcome         0
```

```
dtype: int64
```

```
import numpy as np # Add this line
cols_to_replace = ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']
for col in cols_to_replace:
    df[col] = df[col].replace(0, np.nan)
    df[col].fillna(df[col].median(), inplace=True)
```

<ipython-input-10-7f0fc1778cfe>:5: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values is a copy. For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value, inplace=True) instead.

```
df[col].fillna(df[col].median(), inplace=True)
```

```
from sklearn.preprocessing import StandardScaler
X = df.drop("Outcome", axis=1)
y = df["Outcome"]
X_scaled = StandardScaler().fit_transform(X)
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.3, stratify=y, random_state=42)
```

```
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, confusion_matrix
```

```
# Use estimator instead of base_estimator
```

```

model_bag = BaggingClassifier(estimator=DecisionTreeClassifier(), n_estimators=100)
model_bag.fit(X_train, y_train)
y_pred_bag = model_bag.predict(X_test)

```

```

from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier

```

```

# Use estimator instead of base_estimator if using older version of scikit-learn
model_ada = AdaBoostClassifier(
    DecisionTreeClassifier(max_depth=1), # Pass the base estimator directly
    n_estimators=100,
    learning_rate=1.0
)
model_ada.fit(X_train, y_train)
y_pred_ada = model_ada.predict(X_test)

```

```

import matplotlib.pyplot as plt
errors = [1 - score for score in model_ada.staged_score(X_test, y_test)]
plt.plot(errors)
plt.title("AdaBoost Error Over Iterations")

```

↩ Text(0.5, 1.0, 'AdaBoost Error Over Iterations')



```

from sklearn.ensemble import RandomForestClassifier

```

```

model_rf = RandomForestClassifier(n_estimators=100, max_depth=5, oob_score=True)
model_rf.fit(X_train, y_train)
y_pred_rf = model_rf.predict(X_test)

```

```

from sklearn.ensemble import StackingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier

```

```

base_models = [('svm', SVC(probability=True)), ('knn', KNeighborsClassifier())]
meta_model = LogisticRegression()

```

```

stack_model = StackingClassifier(estimators=base_models, final_estimator=meta_model, cv=5)
stack_model.fit(X_train, y_train)
y_pred_stack = stack_model.predict(X_test)


```

```

ada = AdaBoostClassifier(n_estimators=100)
ada.fit(X_train, y_train)
plt.plot(list(ada.staged_score(X_test, y_test)))

```


↳ [matplotlib.lines.Line2D at 0x79c4a70a7d90]



```
feature_imp = pd.Series(model_rf.feature_importances_, index=X.columns).nlargest(5)
```

```
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import RocCurveDisplay
RocCurveDisplay.from_estimator(model_bag, X_test, y_test)
# Repeat for others
```

↳ <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x79c4a70d1d90>



```
import time
start = time.time()
# Replace 'model' with the actual model you want to use, for example:
model = model_rf # or model_bag, model_ada, stack_model
model.fit(X_train, y_train)
print("Training Time:", time.time() - start)
```

↳ Training Time: 0.40747618675231934

