

```
# -----
# Step 1: Import Libraries
# -----
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.pipeline import Pipeline

# -----
# Question 1: Load Data and Visualize
# -----
# (a) Load dataset
file_path = 'manufacturing.csv' # UPDATE this
manufacturing_df = pd.read_csv(file_path)
print(manufacturing_df.head())
print(manufacturing_df.info())

# (b) Scatter plots
plt.figure(figsize=(14, 5))

plt.subplot(1, 2, 1)
plt.scatter(manufacturing_df['Temperature (°C)'], manufacturing_df['Quality Rating'], alpha=0.7)
plt.xlabel("Temperature (°C)")
plt.ylabel("Quality Rating")
plt.title("Temperature vs Quality Rating")

plt.subplot(1, 2, 2)
plt.scatter(manufacturing_df['Pressure (kPa)'], manufacturing_df['Quality Rating'], alpha=0.7)
plt.xlabel("Pressure (kPa)")
plt.ylabel("Quality Rating")
plt.title("Pressure vs Quality Rating")

plt.tight_layout()
plt.show()
```

```

↩ Temperature (°C) Pressure (kPa) Temperature x Pressure \
0      209.762701      8.050855      1688.769167
1      243.037873      15.812068      3842.931469
2      220.552675      7.843130      1729.823314
3      208.976637      23.786089      4970.736918
4      184.730960      15.797812      2918.345014

```

```

Material Fusion Metric Material Transformation Metric Quality Rating
0      44522.217074      9.229576e+06      99.999971
1      63020.764997      1.435537e+07      99.985703
2      49125.950249      1.072839e+07      99.999758
3      57128.881547      9.125702e+06      99.999975
4      38068.201283      6.303792e+06      100.000000

```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 3957 entries, 0 to 3956
```

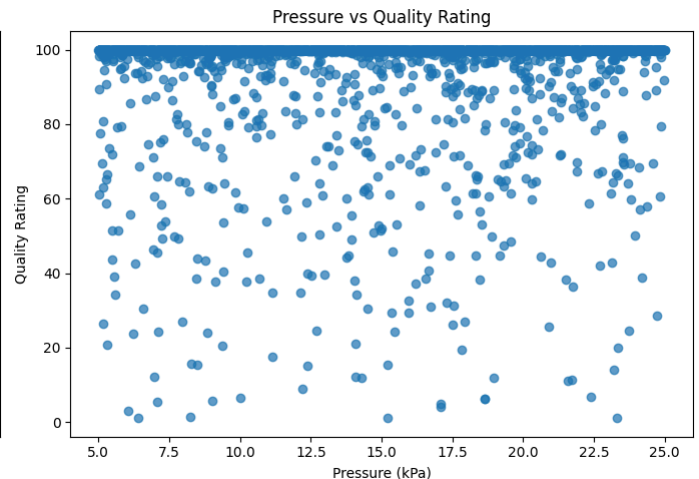
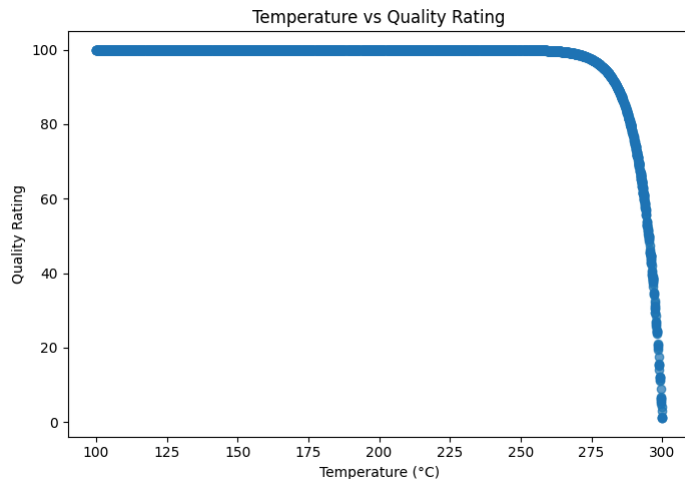
```
Data columns (total 6 columns):
```

#	Column	Non-Null Count	Dtype
0	Temperature (°C)	3957 non-null	float64
1	Pressure (kPa)	3957 non-null	float64
2	Temperature x Pressure	3957 non-null	float64
3	Material Fusion Metric	3957 non-null	float64
4	Material Transformation Metric	3957 non-null	float64
5	Quality Rating	3957 non-null	float64

```
dtypes: float64(6)
```

```
memory usage: 185.6 KB
```

```
None
```



```

# -----
# Question 2: Linear vs Quadratic Model for Temperature
# -----
X = manufacturing_df[['Temperature (°C)']]
y = manufacturing_df['Quality Rating']

# (a) Linear Regression (degree 1)
linear_model = LinearRegression()
linear_model.fit(X, y)
y_pred_linear = linear_model.predict(X)
r2_linear = r2_score(y, y_pred_linear)
mse_linear = mean_squared_error(y, y_pred_linear)
print("Linear Model - R²:", r2_linear, "MSE:", mse_linear)

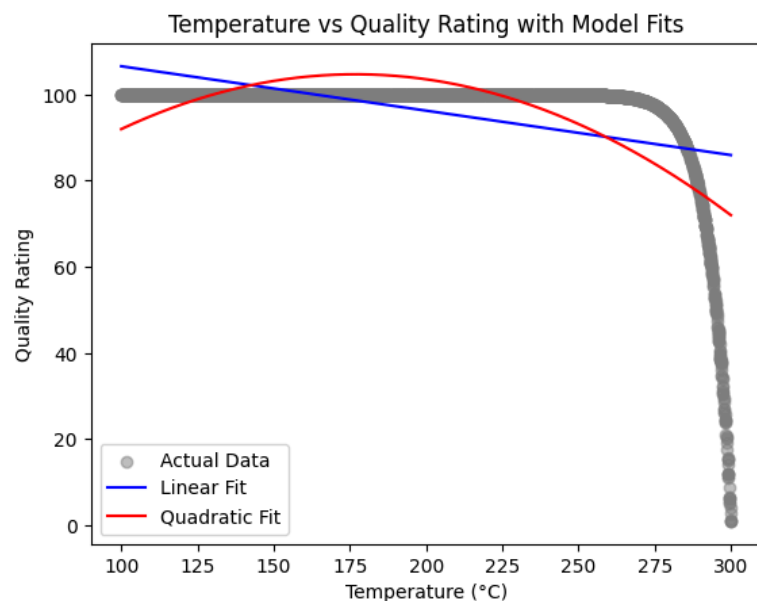
# (b) Quadratic Regression (degree 2)
poly2 = PolynomialFeatures(degree=2)
X_poly2 = poly2.fit_transform(X)
quadratic_model = LinearRegression()
quadratic_model.fit(X_poly2, y)
y_pred_quad = quadratic_model.predict(X_poly2)
r2_quad = r2_score(y, y_pred_quad)
mse_quad = mean_squared_error(y, y_pred_quad)
print("Quadratic Model - R²:", r2_quad, "MSE:", mse_quad)

# (c) Plot both predictions
plt.scatter(X, y, color='gray', alpha=0.5, label='Actual Data')
sorted_idx = X.values.flatten().argsort()

```

```
plt.plot(X.values[sorted_idx], y_pred_linear[sorted_idx], color='blue', label='Linear Fit')
plt.plot(X.values[sorted_idx], y_pred_quad[sorted_idx], color='red', label='Quadratic Fit')
plt.xlabel("Temperature (°C)")
plt.ylabel("Quality Rating")
plt.title("Temperature vs Quality Rating with Model Fits")
plt.legend()
plt.show()
```

Linear Model - R^2 : 0.2127778634271571 MSE: 132.84863630633382
 Quadratic Model - R^2 : 0.4613061729741008 MSE: 90.90793688626003

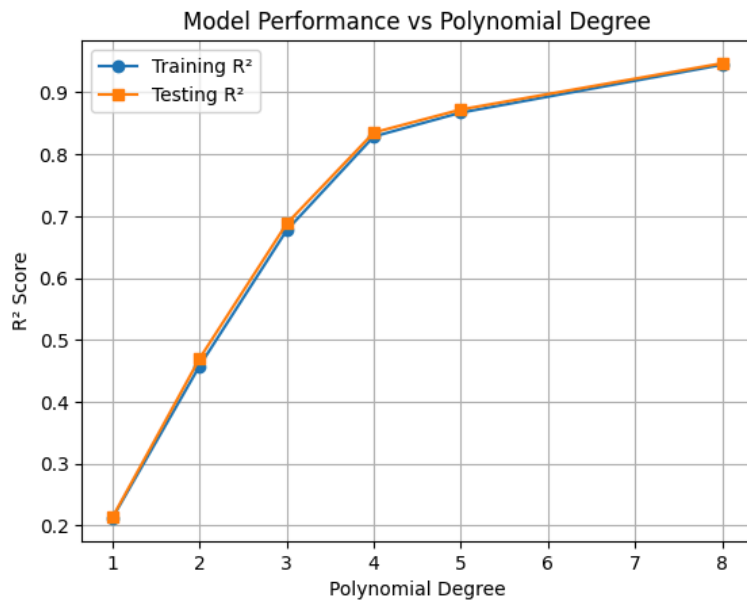


```
# -----
# Question 3: Higher-Degree Polynomials and Overfitting
# -----
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

degrees = [1, 2, 3, 4, 5, 8]
train_r2 = []
test_r2 = []

for d in degrees:
    poly = PolynomialFeatures(degree=d)
    X_train_poly = poly.fit_transform(X_train)
    X_test_poly = poly.transform(X_test)
    model = LinearRegression()
    model.fit(X_train_poly, y_train)
    train_r2.append(model.score(X_train_poly, y_train))
    test_r2.append(model.score(X_test_poly, y_test))


# (c) Plot R² scores
plt.plot(degrees, train_r2, marker='o', label='Training R²')
plt.plot(degrees, test_r2, marker='s', label='Testing R²')
plt.xlabel("Polynomial Degree")
plt.ylabel("R² Score")
plt.title("Model Performance vs Polynomial Degree")
plt.legend()
plt.grid(True)
plt.show()
```



```
# -----
# Question 4: Optimal Degree via Cross-Validation
# -----
cv_degrees = list(range(1, 7))
mean_mse_scores = []

for d in cv_degrees:
    pipeline = Pipeline([
        ('poly', PolynomialFeatures(degree=d)),
        ('model', LinearRegression())
    ])
    neg_mse_scores = cross_val_score(pipeline, X, y, cv=5, scoring='neg_mean_squared_error')
    mean_mse_scores.append(-neg_mse_scores.mean())

# (b) Plot MSE vs degree
plt.plot(cv_degrees, mean_mse_scores, marker='o')
plt.xlabel("Polynomial Degree")
plt.ylabel("Average Cross-Validated MSE")
plt.title("Cross-Validation: Degree vs MSE")
plt.grid(True)
plt.show()
```



```
# -----  
# Question 5: Final Model and Prediction  
# -----  
# Assume best degree (based on CV plot) is 2 or replace with your result  
optimal_degree = cv_degrees[np.argmin(mean_mse_scores)]  
final_poly = PolynomialFeatures(degree=optimal_degree)  
X_final_poly = final_poly.fit_transform(X)  
final_model = LinearRegression()  
final_model.fit(X_final_poly, y)  
  
# (b) Predict for Temperature = 215°C  
temp_val = np.array([[215]])  
temp_poly = final_poly.transform(temp_val)
```