

```

from keras.datasets import imdb
from keras.preprocessing.sequence import pad_sequences
import numpy as np
import matplotlib.pyplot as plt

# Load dataset
(vocab_size, maxlen) = (10000, 200)
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=vocab_size)

# Explore dataset
print("Number of training samples:", len(x_train))
print("Number of test samples:", len(x_test))
print("Average review length:", np.mean([len(i) for i in x_train]))
print("Label distribution:", np.unique(y_train, return_counts=True))

# Pad sequences
x_train = pad_sequences(x_train, maxlen=maxlen)
x_test = pad_sequences(x_test, maxlen=maxlen)

# Visualize padded vs. original
print("Original:", x_train[0][-20:])
print("Padded:", x_train[0])

```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz>  
**17464789/17464789** — 0s 0us/step

```

Number of training samples: 25000
Number of test samples: 25000
Average review length: 238.71364
Label distribution: (array([0, 1]), array([12500, 12500]))
Original: [ 65  16  38 1334  88  12  16 283   5  16 4472 113 103  32
 15  16 5345  19 178  32]
Padded: [  5  25 100  43 838 112  50 670   2   9  35 480 284   5
 150   4 172 112 167   2 336 385  39   4 172 4536 1111  17
 546  38  13 447   4 192  50  16   6 147 2025  19  14  22
  4 1920 4613 469   4  22  71  87  12  16  43 530  38  76
  15  13 1247   4  22  17 515  17  12  16 626  18   2   5
 62 386  12   8 316   8 106   5   4 2223 5244  16 480  66
3785  33   4 130  12  16  38 619   5  25 124  51  36 135
 48  25 1415  33   6  22  12 215  28  77  52   5  14 407
 16  82   2   8   4 107 117 5952  15 256   4   2   7 3766
  5 723  36  71 43 530 476  26 400 317  46   7   4   2
1029  13 104  88   4 381  15 297  98  32 2071  56  26 141
  6 194 7486  18   4 226  22  21 134 476  26 480   5 144
 30 5535  18  51  36  28 224  92  25 104   4 226  65  16
 38 1334  88  12  16 283   5  16 4472 113 103  32  15  16
5345  19 178  32]

```

```

from keras.models import Sequential
from keras.layers import Embedding, SimpleRNN, Dense

model = Sequential([
    Embedding(input_dim=10000, output_dim=32, input_length=200),
    SimpleRNN(32),
    Dense(1, activation='sigmoid')
])

model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['accuracy'])
model.summary()

```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/embedding.py:90: UserWarning: Argument `input\_length` is deprecated. Just warnings.warn(  
Model: "sequential"

Layer (type)	Output Shape	Param #
embedding ( <a href="#">Embedding</a> )	?	0 (unbuilt)
simple_rnn ( <a href="#">SimpleRNN</a> )	?	0 (unbuilt)
dense ( <a href="#">Dense</a> )	?	0 (unbuilt)

Total params: 0 (0.00 B)  
Trainable params: 0 (0.00 B)  
Non-trainable params: 0 (0.00 B)

```
history = model.fit(x_train, y_train, epochs=5, batch_size=128, validation_split=0.2)
```

```
# Plot accuracy and loss
```


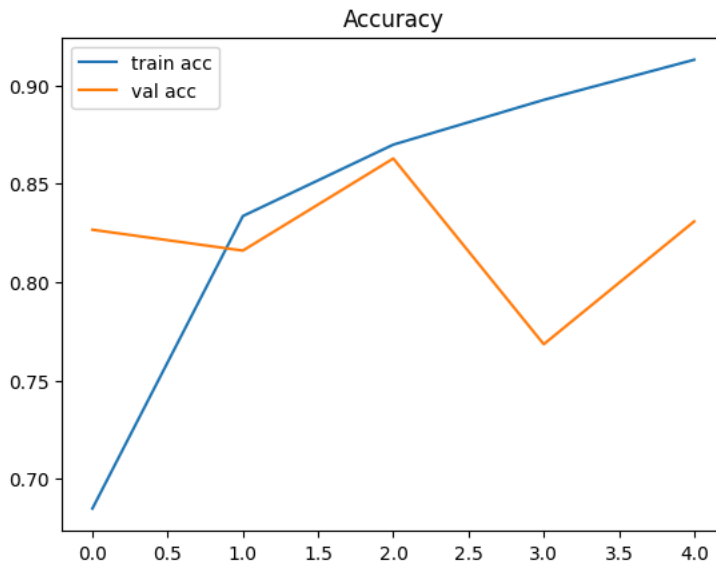
```
plt.plot(history.history['accuracy'], label='train acc')
plt.plot(history.history['val_accuracy'], label='val acc')
plt.title('Accuracy')
plt.legend()
plt.show()
```

```
plt.plot(history.history['loss'], label='train loss')
plt.plot(history.history['val_loss'], label='val loss')
plt.title('Loss')
plt.legend()
plt.show()
```

```
# Evaluate
```

```
test_loss, test_acc = model.evaluate(x_test, y_test)
print("Test Accuracy:", test_acc)
```

```
Epoch 1/5
157/157 — 13s 71ms/step - accuracy: 0.5949 - loss: 0.6456 - val_accuracy: 0.8266 - val_loss: 0.4102
Epoch 2/5
157/157 — 19s 59ms/step - accuracy: 0.8201 - loss: 0.4093 - val_accuracy: 0.8160 - val_loss: 0.4006
Epoch 3/5
157/157 — 10s 55ms/step - accuracy: 0.8706 - loss: 0.3207 - val_accuracy: 0.8628 - val_loss: 0.3399
Epoch 4/5
157/157 — 11s 61ms/step - accuracy: 0.8930 - loss: 0.2710 - val_accuracy: 0.7684 - val_loss: 0.4727
Epoch 5/5
157/157 — 11s 64ms/step - accuracy: 0.9095 - loss: 0.2331 - val_accuracy: 0.8308 - val_loss: 0.3813
```



```
782/782 — 9s 11ms/step - accuracy: 0.8371 - loss: 0.3845
Test Accuracy: 0.8365600109100342
```

```
SimpleRNN(64)
```

```
↗ <SimpleRNN name=simple_rnn_1, built=False>
```

```
from keras.layers import LSTM
```

```
model = Sequential([
    Embedding(input_dim=10000, output_dim=32, input_length=200),
    LSTM(32),
    Dense(1, activation='sigmoid')
])
```

```
# Change vocab_size to 5000 or 2000 in imdb.load_data
from keras.layers import Dropout
```

```
model = Sequential([
    Embedding(input_dim=10000, output_dim=32, input_length=200),
    SimpleRNN(32, dropout=0.2),
    Dense(1, activation='sigmoid')
])
```